

Machine Learning & Data Mining

Classification

Kyung-Ah Sohn

Ajou University

Content

- Introduction to supervised approach
 - Classification
 - Regression
 - Model selection
- Classification algorithms
 - KNN
 - SVM
 - Decision tree
 - Random Forest, Ensemble approach

INTRODUCTION TO SUPERVISED LEARNING


Supervised Learning

Regression



vs

Classification



Supervised learning approach

Feature Space \mathcal{X}

Words in a document



Label Space \mathcal{Y}

"Sports"
"News"
"Science"
...

Discrete Labels
Classification




Share Price
"\$ 24.50"

Continuous Labels
Regression


Task: Given $X \in \mathcal{X}$, predict $Y \in \mathcal{Y}$.

Classification: Training vs. Test

1) Learn from training data



2) Map unseen (new) data




Explain whether each scenario is a classification or regression problem, and provide n (# of samples) and p (# of features).

- (a) We collect a set of data on the top 500 firms in the US. For each firm we record profit, number of employees, industry and the CEO salary. We are interested in understanding which factors affect CEO salary.
- (b) We are considering launching a new product and wish to know whether it will be a *success* or a *failure*. We collect data on 20 similar products that were previously launched. For each product we have recorded whether it was a success or failure, price charged for the product, marketing budget, competition price, and ten other variables.
- (c) We are interesting in predicting the % change in the US dollar in relation to the weekly changes in the world stock markets. Hence we collect weekly data for all of 2012. For each week we record the % change in the dollar, the % change in the US market, the % change in the British market, and the % change in the German market.

Classification

- **Learning:** Induce classifiers from training data



- **Prediction:** use the learned classifier to predict labels for new data



Training vs. Test data

- **Training data**: to learn a classifier (using known labels)

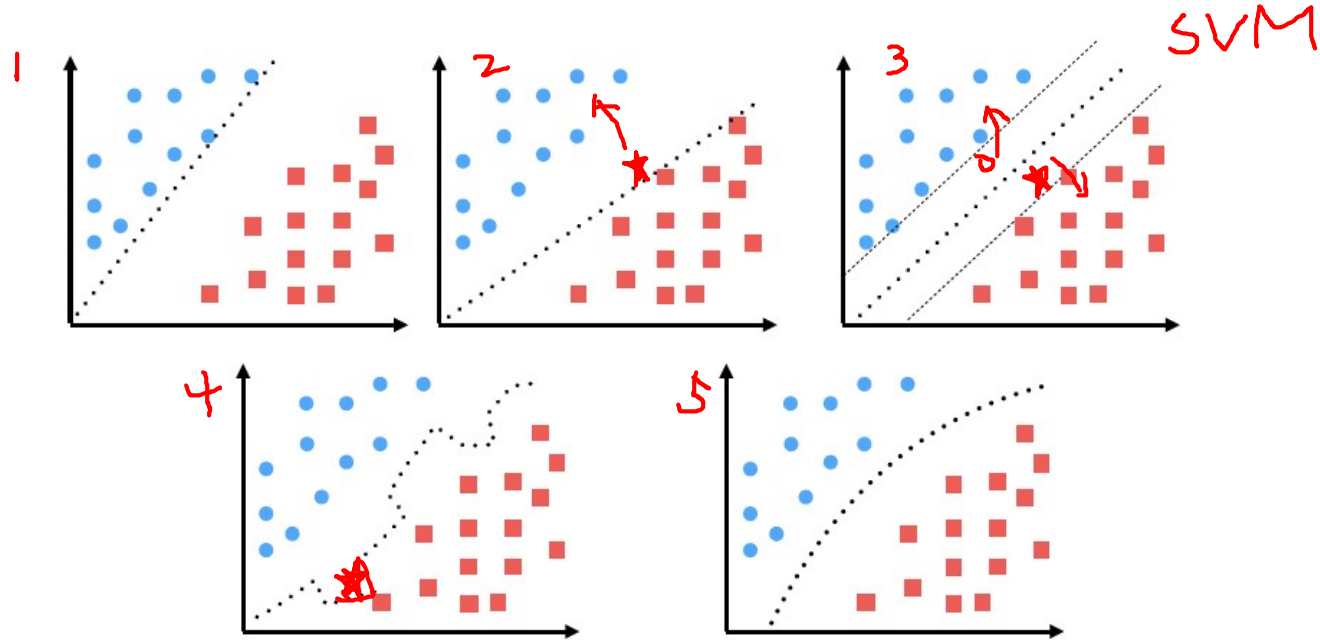
| | team | coach | play | ball | score | game | win | goal | label |
|-------|------|-------|------|------|-------|------|-----|------|-------|
| Doc_1 | 3 | 9 | 0 | 2 | 5 | 7 | 0 | 0 | 1 |
| Doc_2 | 0 | 7 | 5 | 4 | 0 | 1 | 2 | 1 | 0 |
| ... | | | | ... | | | | | ... |
| Doc_n | 0 | 1 | 0 | 0 | 1 | 2 | 2 | 0 | 1 |

- **Test data**: to predict a label or to evaluate the classifier


| | team | coach | play | ball | score | game | win | goal | label |
|--------------------|------|-------|------|------|-------|------|-----|------|-------|
| Doc_new | 3 | 9 | 0 | 2 | 5 | 7 | 0 | 0 | ? |

Model selection

- Which one is better?




❖ Over-fitting for training data




Is red boundary is better than blue one?

❖ Over-fitting for training data

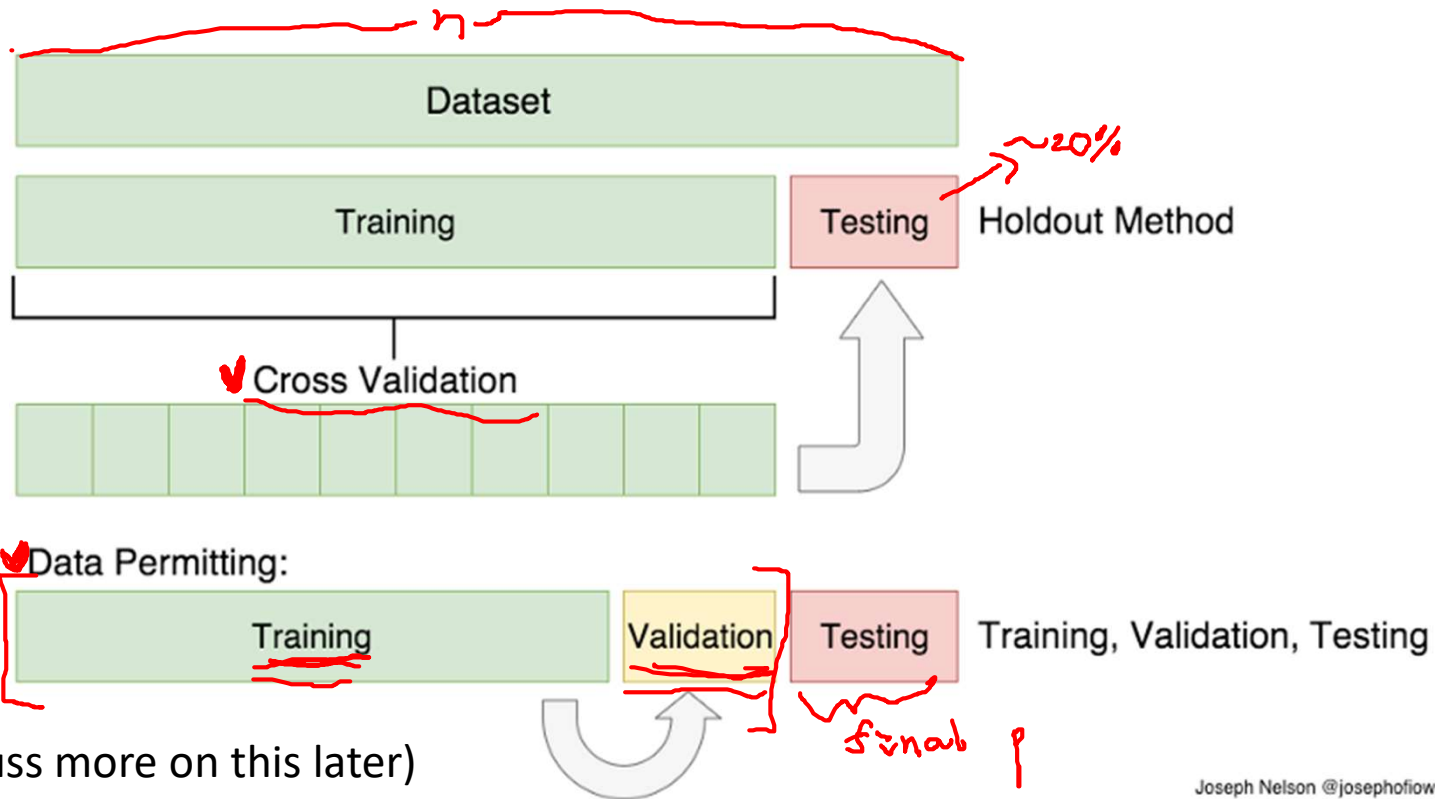


Train/Test split

- The performance of the classifier is measured on the test set
- What if the test set is the same as the training set?
 - 'training error', 'training accuracy' vs. 'test error', 'test accuracy'
- Is the training error always zero?
- Why do we need a separate test set?



Train/Validation/Test set



Joseph Nelson @josephofiowa

Classification algorithms

- K-Nearest Neighbor
- Support Vector Machine (SVM)
- Decision tree, Random Forest
- Naïve Bayes
- Logistic regression
- Neural networks. DL
- Bayesian networks
- ...

K-NEAREST NEIGHBOR ALGORITHM

Classification example

Training data

| No. | v_1 =Strength | v_2 =Smooth | y =Quality |
|-------|-----------------|---------------|--------------|
| x_1 | 5 | 5 | Good |
| x_2 | 1 | 2 | Bad |
| x_3 | 6 | 7 | Good |
| x_4 | 3 | 5 | Bad |

Test data

| | | | |
|-------|---|---|---|
| x_5 | 2 | 3 | ? |
|-------|---|---|---|

k-Nearest Neighbors (k-NN) classification

- Determine k and the distance metric to use
- For a given test sample
 1. Calculate distances of all training samples to the test sample
 2. Pick k closest training samples
 3. Calculate majority

Classification example

- Compute the distance between test data and all the training samples

| No. | v_1 =Strength | v_2 =Smooth | y =Quality |
|-------------------------|-----------------|---------------|--------------|
| x_1 | <u>5</u> | <u>5</u> | Good |
| <u>x_2</u> | 1 | 2 | <u>Bad</u> |
| x_3 | 6 | 7 | Good |
| x_4 | 3 | 5 | <u>Bad</u> |
| x_5 | <u>2</u> | <u>3</u> | ? |


$$d(\mathbf{x}_1, \mathbf{x}_5) = \sqrt{(5 - 2)^2 + (5 - 3)^2} = \underline{3.6}$$

$$d(\mathbf{x}_2, \mathbf{x}_5) = \sqrt{(1 - 2)^2 + (2 - 3)^2} = \underline{1.4}$$

$$d(\mathbf{x}_3, \mathbf{x}_5) = \sqrt{(6 - 2)^2 + (7 - 3)^2} = 5.6$$

$$d(\mathbf{x}_4, \mathbf{x}_5) = \sqrt{(3 - 2)^2 + (5 - 3)^2} = 2.2$$

Example



Computation for KNN

- What is done in training stage?
- How many computations to predict the label of each sample?

Instance-based learning, lazy learning

K-NN algorithm

- 1-NN
 - Predict the same class label as the *nearest* instance in the training set
- In general, k -NN
 - Find the k closest training points (according to some distance measures, e.g. Euclidean, ...)
 - Predicted class: majority vote of k -neighbors

Neighbor size k

- Choice of k
 - Smaller $k \rightarrow$ higher variance (less stable)
 - Larger $k \rightarrow$ higher bias (less precise)
 - Proper choice of k depends on the dataset
 - Heuristics
 - Cross-validation

Neighbor size – extreme cases


- What is the training error (the error when each training sample is used as a test sample) when $K=1$?
- What happens if $K=n$ (the number of training samples)

k-NN


- No explicit training or model, “lazy learning”
- One of the simplest classification methods
- *Gives the maximum likelihood estimation of the class posterior probabilities*
- Can be used as a baseline model
- Many extensions

Choice of Distance Measure

- Euclidean distance
- Manhattan distance
- L_p norm
- Mahalanobis distance
- 1 - Correlation
- 1 - Cosine similarity
- ...



Manhattan



Euclidean

Data and Application-dependent

Euclidean Distance

- n objects with d attributes
 - $\mathbf{x}_i = (x_1(i), x_2(i), \dots, x_d(i))$, $i=1, \dots, n$
- Most common distance metric is Euclidean distance

$$d_E(i, j) = \left(\sum_{k=1}^d (\underline{x_k(i)} - \underline{x_k(j)})^2 \right)^{\frac{1}{2}}$$

Example

| No. | <u>v_1=Strength</u> | <u>v_2=duration</u> | y=Quality |
|-----|----------------------------------|----------------------------------|-----------|
| 1 | 5 | 5000 | Good |
| 2 | 1 | 4900 | Bad |
| 3 | 6 | 2900 | Good |
| 4 | 3 | 5200 | Bad |
| 5 | 2 | 3300 | ? |

Euclidean distance??

Euclidean distance

- Makes sense in the case where the different measurements are commensurate; each variable measured in the same units
- If the measurements are different, say length in different units, Euclidean distance is not necessarily meaningful
 - Standardization!

Normalization

- Usually assume features should contribute equally.

- Min-max normalization

$$X_{new} = \frac{X - \min(X)}{\max(X) - \min(X)}$$

- Z-score standardization

$$X_{new} = \frac{X - \mu}{\sigma} = \frac{X - \text{Mean}(X)}{\text{StdDev}(X)}$$

Extensions

- L_p metric:


$$dist(i, j) = \left(\sum_{k=1}^d |x_k(i) - x_k(j)|^p \right)^{\frac{1}{p}} \quad \text{where } p \geq 1$$

- Manhattan, city block or L_1 metric:

$$dist(i, j) = \sum_{k=1}^d |x_k(i) - x_k(j)|$$

- L_∞

$$dist(i, j) = \max_k |x_k(i) - x_k(j)|$$



Mahalanobis distance

- The Mahalanobis distance takes into account the correlations of the data set.
- If each of these axes is re-scaled to have unit variance, then the Mahalanobis distance corresponds to the standard Euclidean distance


$$d(X, Y) = \sqrt{(X - Y)^T S^{-1} (X - Y)}$$

S : Covariance matrix

- If the covariance matrix is the identity matrix:

$$d(X, Y) = \sqrt{\sum_{i=1}^m \frac{(x_i - y_i)^2}{s_i^2}}$$

- Which point is closer to the origin (0,0)?




Widely used similarity measures

- Cosine similarity
 - Consider only the direction between the two vectors (not the magnitude)
 - Compute using dot product (inner product)

$$S_{cos}(X, Y) = \cos\theta = \frac{X \cdot Y}{||X|| ||Y||}$$

$$X \cdot Y = x_1y_1 + x_2y_2 + \dots + x_my_m$$


$$||X|| = \sqrt{x_1^2 + x_2^2 + \dots + x_m^2}$$




Cosine similarity

- $S_{\cos}((1,1),(2,2)) =$
- $S_{\cos}((1,1),(3,3)) =$
- $S_{\cos}((1,1),(-1,1)) =$


$$S_{\cos}(X, Y) = \cos\theta = \frac{X \cdot Y}{\|X\| \|Y\|}$$



France and Italy are quite similar
 θ is close to 0°
 $\cos(\theta) \approx 1$



ball and crocodile are not similar
 θ is close to 90°
 $\cos(\theta) \approx 0$



the two vectors are similar but opposite
 the first one encodes (city - country)
 while the second one encodes (country - city)
 θ is close to 180°
 $\cos(\theta) \approx -1$

https://datascience-enthusiast.com/DL/Operations_on_word_vectors.html


(Pearson's) correlation coefficient

- Similarity measure
 - Ranges between -1 and 1
 - measures the strength and direction of the linear relationship between two variables.

$$\rho(\mathbf{x}, \mathbf{y}) = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

$$\bar{x} = \frac{1}{n} \sum_i x_i, \quad \bar{y} = \frac{1}{n} \sum_i y_i$$

$$s_\rho(X, Y) = 0.5 \times (1 - \rho(X, Y))$$



K-NN (Pros)

- + Easy to understand and program
- + Explicit reject option
 - If there is no majority agreement
- + Easy handling of missing values
- + Close to optimal
 - *Asymptotic misclassification rate (as the number of data points goes to infinity) is bounded above by twice the Bayes error rate*

K-NN (Cons)

- *Sensitive to noise, irrelevant features*
- Computationally expensive $O(np)$
- *Large memory requirements*
- More frequent classes dominate result
- Curse of dimensionality
 - “nearest” neighbor may be very far
 - In high dimensions, “nearest” becomes meaningless

Curse of dimensionality

- The size of the data *space* grows exponentially with the number of dimensions
- The size of your data set must also grow exponentially in order to keep the same density...
- In KNN algorithm, it requires two points to be very close on *every axis*
 - Adding a new dimension creates another chance for points to be farther apart
- How to overcome in KNN
 - Add more data...
 - Reduce dimension (feature selection, feature extraction)