


SUPPORT VECTOR MACHINE

Classification

- **Task: given X, predict Y**
 - Labeled data (X & Y)
 - “**Predict** class label (Y)”





Binary classification

SVM - Introduction

- SVMs provide a learning technique for classification
- Solution provided SVM is
 - Theoretically elegant
 - Computationally Efficient
 - Very effective in many large practical problems
- It has a simple geometrical interpretation in a high-dimensional feature space that is nonlinearly related to input space
- By using **kernels** all computations keep simple.

Content

1. Linear SVM
2. Linear SVM: non-separable case
3. Non-linear SVM
 - Kernel trick




History

- The Study on Statistical Learning Theory was started in the 1960s by Vapnik
- Support Vector Machine is a practical learning method based on Statistical Learning Theory
- A simple SVM could beat a sophisticated neural networks with elaborate features in a handwriting recognition task.
- Now deep neural networks could beat SVM

Vladimir Vapnik




Support Vector Machines



Two dimensional
training data


- Find a linear hyperplane (decision boundary) that will separate the data

Support Vector Machines




- One Possible Solution

Support Vector Machines




- Another possible solution

Support Vector Machines




- Other possible solutions

Support Vector Machines



- Which one is better? B_1 or B_2 ?
- How do you define better?


Support Vector Machines



- Find hyperplane **maximizes** the margin $\Rightarrow B_1$ is better than B_2

Linear SVM: Separable case

- Linear decision boundary
- Consider a binary classification problem consisting of N training samples (\mathbf{x}_i, y_i) , $i=1, \dots, N$ where
 - $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{id})^T$: the attributes
 - y_i : class labels (either -1 or 1)
 - Decision boundary: $\mathbf{w} \cdot \mathbf{x} + b = 0$
(or $w_1x_1 + \dots + w_dx_d + b = 0$)




Linear Classifier


$$\begin{aligned} \mathbf{w} \cdot \mathbf{x}_a + b &= 0 \\ \mathbf{w} \cdot \mathbf{x}_b + b &= 0 \\ \Rightarrow \mathbf{w} \cdot (\mathbf{x}_a - \mathbf{x}_b) &= 0 \end{aligned}$$

One possible linear classifier:

$$y = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ -1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b < 0 \end{cases}$$



Linear SVM




Constraints

$$y_i = \begin{cases} 1 & \text{if } w \bullet x_i + b \geq 1 \\ -1 & \text{if } w \bullet x_i + b \leq -1 \end{cases}$$

Margin =

Linear SVM




Constraints:

$$y_i = \begin{cases} 1 & \text{if } w \bullet x_i + b \geq 1 \\ -1 & \text{if } w \bullet x_i + b \leq -1 \end{cases}$$

Objective: **Maximize** Margin = $\frac{2}{\|w\|}$

Learning linear SVM

- **Training phase:** estimate the parameters w and b from the training data
- Objective: maximize margin
- Constraints:
$$\begin{aligned} \mathbf{w} \cdot \mathbf{x}_i + b &\geq 1 \text{ if } y_i = 1 \\ \mathbf{w} \cdot \mathbf{x}_i + b &\leq -1 \text{ if } y_i = -1 \end{aligned}$$

 $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1, i = 1, \dots, N$

Optimization problem

- We want to maximize: $\text{Margin} = \frac{2}{\|\vec{w}\|^2}$
 - Which is equivalent to minimizing: $L(w) = \frac{\|\vec{w}\|^2}{2}$
 - But subjected to the following constraints:

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1, i = 1, \dots, N$$

- This is a constrained (convex) optimization problem
 - Numerical approaches to solve it (e.g., quadratic programming)
 - Lagrange multiplier method:

Linear SVM

- Test phase

- Once the parameters of the decision boundary are found, a test instance z is classified as follows:

$$f(z) = 1 \text{ if } \mathbf{w} \cdot \mathbf{z} + b \geq 0$$

$$f(z) = -1 \text{ otherwise}$$

Example – linear SVM

```
from sklearn.svm import SVC # "Support vector classifier"
```

```
# create SVM classifier
```


```
model = SVC(kernel='linear')
```

```
# train the model using the training data
```

```
model.fit(Xtrain, ytrain)
```

```
# predict the label for test dataset
```


```
Ypred = model.predict(Xtest)
```



<https://jakevdp.github.io/PythonDataScienceHandbook/05.07-support-vector-machines.html>

Linear SVM: non-separable case

- What if the problem is not linearly separable?



Use a soft margin approach


Slack variables for non-separable data

Relax the constraints

$$\mathbf{w} \cdot \mathbf{x}_i + b \geq 1 - \xi \text{ if } y_i = 1$$

$$\mathbf{w} \cdot \mathbf{x}_i + b \leq -1 + \xi \text{ if } y_i = -1$$

$$\text{minimize } L(\mathbf{w}) = \frac{\|\vec{\mathbf{w}}\|^2}{2} + C \left(\sum_{i=1}^N \xi_i \right)^k$$



Linear SVM: non-separable case

- What if the problem is not linearly separable?
 - Introduce slack variables

- Need to minimize:
$$L(w) = \frac{\|\vec{w}\|^2}{2} + C \left(\sum_{i=1}^N \xi_i \right)^k$$

- Subject to:
$$f(\vec{x}_i) = \begin{cases} 1 & \text{if } \vec{w} \bullet \vec{x}_i + b \geq 1 - \xi_i \\ -1 & \text{if } \vec{w} \bullet \vec{x}_i + b \leq -1 + \xi_i \end{cases}$$

- C, k : user-specified parameters

Tuning hyper-parameter


- C for soft margin
 - Large C: margin is hard
 - Smaller C: margin is soft

$$\begin{aligned} \mathbf{w} \cdot \mathbf{x}_i + b &\geq 1 - \xi \text{ if } y_i = 1 \\ \mathbf{w} \cdot \mathbf{x}_i + b &\leq -1 + \xi \text{ if } y_i = -1 \end{aligned}$$

$$\text{minimize } L(\mathbf{w}) = \frac{\|\vec{\mathbf{w}}\|^2}{2} + C \left(\sum_{i=1}^N \xi_i \right)^k$$


```
model = SVC(kernel='linear', C=10.0)
```

C=10.0



```
model = SVC(kernel='linear', C=0.1)
```

C=0.1




Linear SVM - summary

- Training phase: solve a constrained optimization problem to get the decision boundary (w and b)
 - Separable case
 - Non-separable case


- Test phase: given w, b

$$f(z) = 1 \text{ if } \mathbf{w} \cdot \mathbf{z} + b \geq 0$$

$$f(z) = -1 \text{ otherwise}$$



Problems with linear SVM



What if the decision function is not a linear?

Non-linear SVMs


- Datasets that are linearly separable with some noise work out great:



- But what are we going to do if the dataset is just too hard?




- How about... mapping data to a higher-dimensional space:




Nonlinear Support Vector Machines

- What if decision boundary is not linear?




Nonlinear SVM

- Transform data into a new space so that a linear boundary can be used to separate the instance



$$\mathbf{x} \rightarrow \Phi(\mathbf{x})$$



SVM with polynomial kernel visualization

<http://www.youtube.com/watch?v=3liCbRZPrZA>




Issues in nonlinear SVM

- It is not clear what type of mapping function to use
- Solving optimization problems in a high-dimensional space can be computationally expensive

➔ Use **Kernel Trick**

Extension to Non-linear Decision Boundary

- Possible problem of the transformation
 - High computation burden and hard to get a good estimate
- SVM solves these two issues simultaneously
 - [Kernel tricks](#) for efficient computation
 - Minimizing $\|\mathbf{w}\|^2$ can lead to a “good” classifier



Learning nonlinear SVM

- Training phase:

- Need to minimize: $L(w) = \frac{\|w\|^2}{2}$

- subject to $y_i(\mathbf{w} \cdot \Phi(\mathbf{x}_i) + b) \geq 1, i = 1, \dots, N$

- Test phase: for a new instance z ,

$$f(z) = 1 \text{ if } \mathbf{w} \cdot \Phi(\mathbf{z}) + b \geq 0$$

$$f(z) = -1 \text{ otherwise}$$

Kernel trick

$$\mathbf{w} \cdot \Phi(\mathbf{z}) + b = \sum_{i=1}^n \lambda_i y_i \Phi(x_i) \cdot \Phi(\mathbf{z}) + b$$

Dot product in a high-dimensional space...
A kind of similarity measure

- There may exist a kernel function K such that

$$K(\mathbf{u}, \mathbf{v}) = \Phi(\mathbf{u}) \Phi(\mathbf{v})$$

- The dot product in a the transformed space can be expressed in terms of a similarity in the original space
- e.g. $K(\mathbf{u}, \mathbf{v}) = (\mathbf{u} \cdot \mathbf{v} + 1)^2$ for $\Phi(\mathbf{u}) = (u_1^2, u_2^2, \sqrt{2}u_1, \sqrt{2}u_2, 1)$

Example Transformation

- Define the kernel function $K(\mathbf{x}, \mathbf{y})$ as

$$K(\mathbf{x}, \mathbf{y}) = (1 + x_1y_1 + x_2y_2)^2$$

- Consider the following transformation

$$\phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right) = (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1x_2)$$

$$\phi\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix}\right) = (1, \sqrt{2}y_1, \sqrt{2}y_2, y_1^2, y_2^2, \sqrt{2}y_1y_2)$$

$$\begin{aligned} \langle \phi\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right), \phi\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix}\right) \rangle &= (1 + x_1y_1 + x_2y_2)^2 \\ &= K(\mathbf{x}, \mathbf{y}) \end{aligned}$$

- The inner product can be computed by K without going through the map $\phi(\cdot)$

Kernel Trick

- K: kernel function
 - The kernel functions can be expressed as the dot product between two input vectors in some high-dimensional space
 - Computing the dot product using kernel functions is considerably cheaper than using the transformed attribute
 - We do not have to know the exact form of the mapping function Φ

Examples of Kernel Functions

- Polynomial kernel with degree d

$$K(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y} + 1)^d$$

- Radial basis function (RBF) kernel with width σ

$$K(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / (2\sigma^2))$$

- Sigmoid with parameter κ and θ

$$K(\mathbf{x}, \mathbf{y}) = \tanh(\kappa \mathbf{x}^T \mathbf{y} + \theta)$$

Examples of Kernel Functions

- Polynomial kernel with degree d

$$K(\mathbf{x}, \mathbf{y}) = (\mathbf{x}^T \mathbf{y} + 1)^d$$

- Radial basis function kernel with width σ

$$K(\mathbf{x}, \mathbf{y}) = \exp(-\|\mathbf{x} - \mathbf{y}\|^2 / (2\sigma^2))$$


–Closely related to radial basis function neural networks

- Sigmoid with parameter κ and θ


$$K(\mathbf{x}, \mathbf{y}) = \tanh(\kappa \mathbf{x}^T \mathbf{y} + \theta)$$

Example - RBF

```
clf = SVC(kernel='linear')  
clf.fit(X, y)
```




```
clf = SVC(kernel='rbf', C=1E6)  
clf.fit(X, y)
```







<https://jakevdp.github.io/PythonDataScienceHandbook/05.07-support-vector-machines.html>

Example



Non-linearly separable case



<https://github.com/rasbt/python-machine-learning-book/blob/master/code/ch03/ch03.ipynb>

SVM - summary

- Convex optimization problem in which efficient algorithms are available
- **Maximizing margin** of the decision boundary
- **Attribute transformation** to a high-dimensional space and **kernel trick**
- The user must still provide other parameters such as the type of kernel function and the cost function C for slack variables
- For binary classification. Can be extended to multi-class problems

Classification - Summary

- Classification algorithms discussed so far
 - KNN
 - Decision Tree
 - Random Forest
 - Support Vector Machine

Things to know

- What is done in training and test phase, respectively
- Time complexity
- Decision boundary
- Model parameters
- Hyper-parameters